

# Fatal Fungi: Enhancing Mushroom Safety through Backpropagated Multi-Layer Perceptron Neural Network Identification

## CSS 581 - Final Project ML Implementation

### Yasmine Subbagh

#### Importing the Data

The dataset used to train and test this neural network is a mixture of two free open-source Kaggle datasets. This method was chosen in order to maximize the training dataset size to create a more accurate classifier.

```
In [1]: import kagglehub
import numpy as np # linear algebra
np.random.seed(10)
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
pd.set_option('display.max_columns', None)
import matplotlib.pyplot as plt
import seaborn as sns
import os
import warnings

# Required magic to display matplotlib plots in notebooks
%matplotlib inline

from sklearn import metrics
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import PCA
```

```
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
```

```
from .autonotebook import tqdm as notebook_tqdm
```

```
In [2]: # download from kaggle
path1 = kagglehub.dataset_download("vishalpnaik/mushroom-classification-edition")
```

```
path2 = kagglehub.dataset_download("ulrikthgepedersen/mushroom-attributes")

files1 = os.listdir(path1)
files2 = os.listdir(path2)

file1 = os.path.join(path1, "mushroom.csv")
file2 = os.path.join(path2, "mushroom.csv")

# Read the file into a list and clean up the byte-like characters
with open(file2, "r") as file:
    lines = file.readlines()

# Clean the lines
cleaned_lines = [line.replace("b'", "").replace("'", "").strip() for line in lines]

# Create a DataFrame
header = cleaned_lines[0].split(",")
data = [row.split(",") for row in cleaned_lines[1:]]

# Load the datasets
df1 = pd.read_csv(file1)
df2 = pd.DataFrame(data, columns=header)

# Print the size of the datasets
print(f"Dataset 1 Size: {df1.shape[0]} rows, {df1.shape[1]} columns")
print(f"Dataset 2 Size: {df2.shape[0]} rows, {df2.shape[1]} columns")
```

Dataset 1 Size: 61069 rows, 21 columns

Dataset 2 Size: 8124 rows, 23 columns

Next, the data needs to be cleaned up and combined into one dataset.

```
In [3]: # Display the DataFrame
print(df2.head())
```

	cap-shape	cap-surface	cap-color	bruises%3F	odor	gill-attachment	\
0	x	s	n	t	p		f
1	x	s	y	t	a		f
2		s	w	t	l		f
3	x	y	w	t	p		f
4	x	s	g	f	n		f

	gill-spacing	gill-size	gill-color	stalk-shape	stalk-root	\
0	c	n	k	e	e	
1	c		k	e	c	
2	c		n	e	c	
3	c	n	n	e	e	
4	w		k	t	e	

	stalk-surface-above-ring	stalk-surface-below-ring	stalk-color-above-ring	\
0		s	s	w
1		s	s	w
2		s	s	w
3		s	s	w
4		s	s	w

	stalk-color-below-ring	veil-type	veil-color	ring-number	ring-type	\
0		w	p	w	o	p
1		w	p	w	o	p
2		w	p	w	o	p
3		w	p	w	o	p
4		w	p	w	o	e

	spore-print-color	population	habitat	class
0	k	s	u	p
1	n	n	g	e
2	n	n	m	e
3	k	s	u	p
4	n	a	g	e

```
In [4]: # Display the DataFrame
print(df1.head())
```

	class	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed
0	p	15.26	x	g	o	f
1	p	16.60	x	g	o	f
2	p	14.07	x	g	o	f
3	p	14.17	f	h	e	f
4	p	14.64	x	h	o	f

	gill-attachment	gill-spacing	gill-color	stem-height	stem-width	stem-root
0	e	NaN	w	16.95	17.09	s
1	e	NaN	w	17.99	18.19	s
2	e	NaN	w	17.80	17.74	s
3	e	NaN	w	15.77	15.98	s
4	e	NaN	w	16.53	17.20	s

	stem-surface	stem-color	veil-type	veil-color	has-ring	ring-type
0	y	w	u	w	t	g
1	y	w	u	w	t	g
2	y	w	u	w	t	g
3	y	w	u	w	t	p
4	y	w	u	w	t	p

	spore-print-color	habitat	season
0	NaN	d	w
1	NaN	d	u
2	NaN	d	w
3	NaN	d	w
4	NaN	d	w

After exploration of the dataset size and attributes, the executive decision to not to continue to use dataset 2 moving forward. This is because it is significantly smaller and the attributes do not align closely, lots of attributes would need to be cut to combine the dataset.

Next, we need to identify duplicate rows and delete them. As well as replace missing categorical data with a "unknown" value to avoid deleting too many data rows. But rows with missing numerical data can be deleted.

```
In [5]: # drop duplicates
print("Number of duplicate rows:", df1.duplicated().sum())
df1 = df1.drop_duplicates()

# drop rows with missing numerical values
numerical_columns = df1.select_dtypes(include=['int64', 'float64']).columns
print(f"Missing values in numerical columns:\n{df1[numerical_columns].isnull()}")
df1 = df1.dropna(subset=numerical_columns)
```

```
# replace missing categorical values with 'Unknown'
categorical_columns = df1.select_dtypes(include=['object']).columns
print(f"Missing values in categorical columns:\n{df1[categorical_columns].isna()}")
df1[categorical_columns] = df1[categorical_columns].fillna('Unknown')
```

Number of duplicate rows: 146

Missing values in numerical columns:

```
cap-diameter    0
stem-height     0
stem-width      0
```

dtype: int64

Missing values in categorical columns:

```
class                0
cap-shape            0
cap-surface         14120
cap-color            0
does-bruise-or-bleed 0
gill-attachment     9855
gill-spacing        25062
gill-color           0
stem-root           51536
stem-surface        38122
stem-color           0
veil-type           57746
veil-color          53510
has-ring            0
ring-type           2471
spore-print-color    54597
habitat              0
season              0
```

dtype: int64

We will not be dropping any columns as it is important for the model to be able to take into account as many attributes to be able to most accurately classify the species, every detail is important and small nuances can indicate different species (which could change class type). Additionally, feature creating is not possible for this type of dataset.

Next, we will rename and reformat the columns and data to be hot coded or scaled to be able to be used in the perceptron layers with accuracy.

```
In [6]: # Modify 'class' column to 'edible' with one-hot encoding (0 for 'p' and 1 for 'e')
df1['edible'] = df1['class'].map({'p': 0, 'e': 1})

# One-hot encode the 'does-bruise-or-bleed' column
df1['does_bruise_or_bleed'] = df1['does-bruise-or-bleed'].map({'t': 1, 'f': 0})

# One-hot encode the 'has-ring' column
df1['has-ring'] = df1['has-ring'].map({'t': 1, 'f': 0})
```

```
# scale the cap diamter, stem height, and stem width to be within 0 and 1
scaler = MinMaxScaler(feature_range=(0, 1))
df1['cap-diameter'] = scaler.fit_transform(df1[['cap-diameter']])
df1['stem-height'] = scaler.fit_transform(df1[['stem-height']])
df1['stem-width'] = scaler.fit_transform(df1[['stem-width']])

# Use pd.get_dummies() to one-hot encode the 'gill-attachment' column
df_encoded_gillAttachment = pd.get_dummies(df1['gill-attachment'], prefix='g')
df_encoded_gillAttachment = df_encoded_gillAttachment * 1
df1 = pd.concat([df1, df_encoded_gillAttachment], axis=1)

# Use pd.get_dummies() to one-hot encode the 'cap-shape' column
df_encoded_cap_shape = pd.get_dummies(df1['cap-shape'], prefix='cap_shape',
df1 = pd.concat([df1, df_encoded_cap_shape], axis=1)

# Use pd.get_dummies() to one-hot encode the 'cap-surface' column
df_encoded_cap_surface = pd.get_dummies(df1['cap-surface'], prefix='cap_surf
df1 = pd.concat([df1, df_encoded_cap_surface], axis=1)

# Use pd.get_dummies() to one-hot encode the 'cap-color' column
df_encoded_cap_color = pd.get_dummies(df1['cap-color'], prefix='cap_color',
df1 = pd.concat([df1, df_encoded_cap_color], axis=1)

# Use pd.get_dummies() to one-hot encode the 'gill-spacing' column
df_encoded_gill_spacing = pd.get_dummies(df1['gill-spacing'], prefix='gill_s
df1 = pd.concat([df1, df_encoded_gill_spacing], axis=1)

# Use pd.get_dummies() to one-hot encode the 'gill-color' column
df_encoded_gill_color = pd.get_dummies(df1['gill-color'], prefix='gill_color
df1 = pd.concat([df1, df_encoded_gill_color], axis=1)

# Use pd.get_dummies() to one-hot encode the 'stem-root' column
df_encoded_stem_root = pd.get_dummies(df1['stem-root'], prefix='stem_root',
df1 = pd.concat([df1, df_encoded_stem_root], axis=1)

# Use pd.get_dummies() to one-hot encode the 'stem-surface' column
df_encoded_stem_surface = pd.get_dummies(df1['stem-surface'], prefix='stem_s
df1 = pd.concat([df1, df_encoded_stem_surface], axis=1)

# Use pd.get_dummies() to one-hot encode the 'stem-color' column
df_encoded_stem_color = pd.get_dummies(df1['stem-color'], prefix='stem_color
df1 = pd.concat([df1, df_encoded_stem_color], axis=1)

# Use pd.get_dummies() to one-hot encode the 'veil-color' column
df_encoded_veil_color = pd.get_dummies(df1['veil-color'], prefix='veil_color
df1 = pd.concat([df1, df_encoded_veil_color], axis=1)

# Use pd.get_dummies() to one-hot encode the 'ring-type' column
df_encoded_ring_type = pd.get_dummies(df1['ring-type'], prefix='ring-type',
df1 = pd.concat([df1, df_encoded_ring_type], axis=1)
```

```

# Use pd.get_dummies() to one-hot encode the 'spore-print-color' column
df_encoded_spore_print_color = pd.get_dummies(df1['spore-print-color'], prefix='spore-print-color')
df1 = pd.concat([df1, df_encoded_spore_print_color], axis=1)

# Use pd.get_dummies() to one-hot encode the 'season' column
df_encoded_season = pd.get_dummies(df1['season'], prefix='season', drop_first=True)
df1 = pd.concat([df1, df_encoded_season], axis=1)

# Use pd.get_dummies() to one-hot encode the 'habitat' column
df_encoded_habitat = pd.get_dummies(df1['habitat'], prefix='habitat', drop_first=True)
df1 = pd.concat([df1, df_encoded_habitat], axis=1)

#only one value for veil-type, drop column
df1 = df1.drop(columns=['veil-type'])

# Drop the old columns
df1 = df1.drop(columns=['habitat', 'season', 'spore-print-color', 'spore-print-color'])
print(df1.head())

```

	cap-diameter	stem-height	stem-width	has-ring	edible	\
0	0.240155	0.499705	0.164469	1	0	
1	0.261782	0.530366	0.175055	1	0	
2	0.220949	0.524764	0.170725	1	0	
3	0.222563	0.464917	0.153787	1	0	
4	0.230148	0.487323	0.165528	1	0	

	does_bruise_or_bleed	gill_attachment_a	gill_attachment_d	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	gill_attachment_e	gill_attachment_f	gill_attachment_p	gill_attachment_s	\
0	1	0	0		
1	1	0	0		
2	1	0	0		
3	1	0	0		
4	1	0	0		

	gill_attachment_x	cap_shape_b	cap_shape_c	cap_shape_f	cap_shape_o	\
0	0	0	0	0	0	
1	0	0	0	0	0	

2	0	0	0	0	0
3	0	0	0	1	0
4	0	0	0	0	0

	cap_shape_p	cap_shape_s	cap_shape_x	cap_surface_d	cap_surface_e	\
0	0	0	1	0	0	
1	0	0	1	0	0	
2	0	0	1	0	0	
3	0	0	0	0	0	
4	0	0	1	0	0	

	cap_surface_g	cap_surface_h	cap_surface_i	cap_surface_k	cap_surface_l	\
0	1	0	0	0	0	
1	1	0	0	0	0	
2	1	0	0	0	0	
3	0	1	0	0	0	
4	0	1	0	0	0	

	cap_surface_s	cap_surface_t	cap_surface_w	cap_surface_y	cap_color_b	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	cap_color_e	cap_color_g	cap_color_k	cap_color_l	cap_color_n	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	1	0	0	0	0	
4	0	0	0	0	0	

	cap_color_o	cap_color_p	cap_color_r	cap_color_u	cap_color_w	\
0	1	0	0	0	0	
1	1	0	0	0	0	
2	1	0	0	0	0	
3	0	0	0	0	0	
4	1	0	0	0	0	

	cap_color_y	gill_spacing_c	gill_spacing_d	gill_spacing_f	gill_color_b	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	gill_color_e	gill_color_f	gill_color_g	gill_color_k	gill_color_n	\
--	--------------	--------------	--------------	--------------	--------------	---



0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

	gill_color_o	gill_color_p	gill_color_r	gill_color_u	gill_color_w	\
0	0	0	0	0	1	
1	0	0	0	0	1	
2	0	0	0	0	1	
3	0	0	0	0	1	
4	0	0	0	0	1	

	gill_color_y	stem_root_b	stem_root_c	stem_root_f	stem_root_r	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	stem_root_s	stem_surface_f	stem_surface_g	stem_surface_h	\
0	1	0	0	0	
1	1	0	0	0	
2	1	0	0	0	
3	1	0	0	0	
4	1	0	0	0	

	stem_surface_i	stem_surface_k	stem_surface_s	stem_surface_t	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	stem_surface_y	stem_color_b	stem_color_e	stem_color_f	stem_color_g	\
0	1	0	0	0	0	
1	1	0	0	0	0	
2	1	0	0	0	0	
3	1	0	0	0	0	
4	1	0	0	0	0	

	stem_color_k	stem_color_l	stem_color_n	stem_color_o	stem_color_p	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	stem_color_r	stem_color_u	stem_color_w	stem_color_y	veil_color_e	\
0	0	0	1	0	0	

1	0	0	1	0	0
2	0	0	1	0	0
3	0	0	1	0	0
4	0	0	1	0	0

	veil_color_k	veil_color_n	veil_color_u	veil_color_w	veil_color_y	\
0	0	0	0	1	0	
1	0	0	0	1	0	
2	0	0	0	1	0	
3	0	0	0	1	0	
4	0	0	0	1	0	

	ring-type_e	ring-type_f	ring-type_g	ring-type_l	ring-type_m	\
0	0	0	1	0	0	
1	0	0	1	0	0	
2	0	0	1	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	ring-type_p	ring-type_r	ring-type_z	spore_print_color_g	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	1	0	0	0	
4	1	0	0	0	

	spore_print_color_k	spore_print_color_n	spore_print_color_p	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	spore_print_color_r	spore_print_color_u	spore_print_color_w	season_a	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	


	season_s	season_u	season_w	habitat_d	habitat_g	habitat_h	habitat_l	\
0	0	0	1	1	0	0	0	
1	0	1	0	1	0	0	0	
2	0	0	1	1	0	0	0	
3	0	0	1	1	0	0	0	
4	0	0	1	1	0	0	0	

	habitat_m	habitat_p	habitat_u	habitat_w
--	-----------	-----------	-----------	-----------


0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

Now that we have finished up with the data pre-processing, we can do some quick data exploration.

```
In [7]: # Pairplot of a subset of features (you can select a few important ones to n
subset = df1[['stem-width', 'stem-height', 'edible']]
sns.pairplot(subset, hue='edible') # Color by 'edible' class
plt.show()
```



```
In [8]: # Pairplot of a subset of features (you can select a few important ones to n
subset = df1[['cap-diameter', 'has-ring', 'edible']]
sns.pairplot(subset, hue='edible') # Color by 'edible' class
plt.show()
```



As we can see from the pairplots above, there seems to be a strong grouping features to the edible (or not) classification. This indicates that a model should be able to classify mushrooms given the data

```
In [9]: # Train a RandomForest classifier
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(df1.drop(columns=['edible']), df1['edible'])

# Get feature importances
importances = model.feature_importances_

feature_importance_df = pd.DataFrame({
    'Feature': df1.drop(columns=['edible']).columns,
    'Importance': importances
})

feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)

# List the top 5 most important features
```


```

top_5_features = feature_importance_df.head(5)
print(top_5_features)

# Plot
plt.figure(figsize=(10, 6))
plt.barh(feature_importance_df['Feature'], feature_importance_df['Importance'])
plt.xlabel('Feature Importance')
plt.title('Feature Importance from Random Forest')
plt.show()

```

	Feature	Importance
2	stem-width	0.079685
1	stem-height	0.055777
0	cap-diameter	0.054871
81	stem_color_w	0.034807
43	gill_spacing_d	0.028843



As we can see from the feature importance graph above, the stem-width has strong feature importance compared to the rest of the features. However, since the data is hot coded, the other features could not be truly represented.

```

In [10]: # Separate features and target
X = df1.drop(columns=['edible'])
y = df1['edible']

# Perform PCA and reduce to 3 components
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
pca = PCA(n_components=3)

```

```
X_pca = pca.fit_transform(X_scaled)

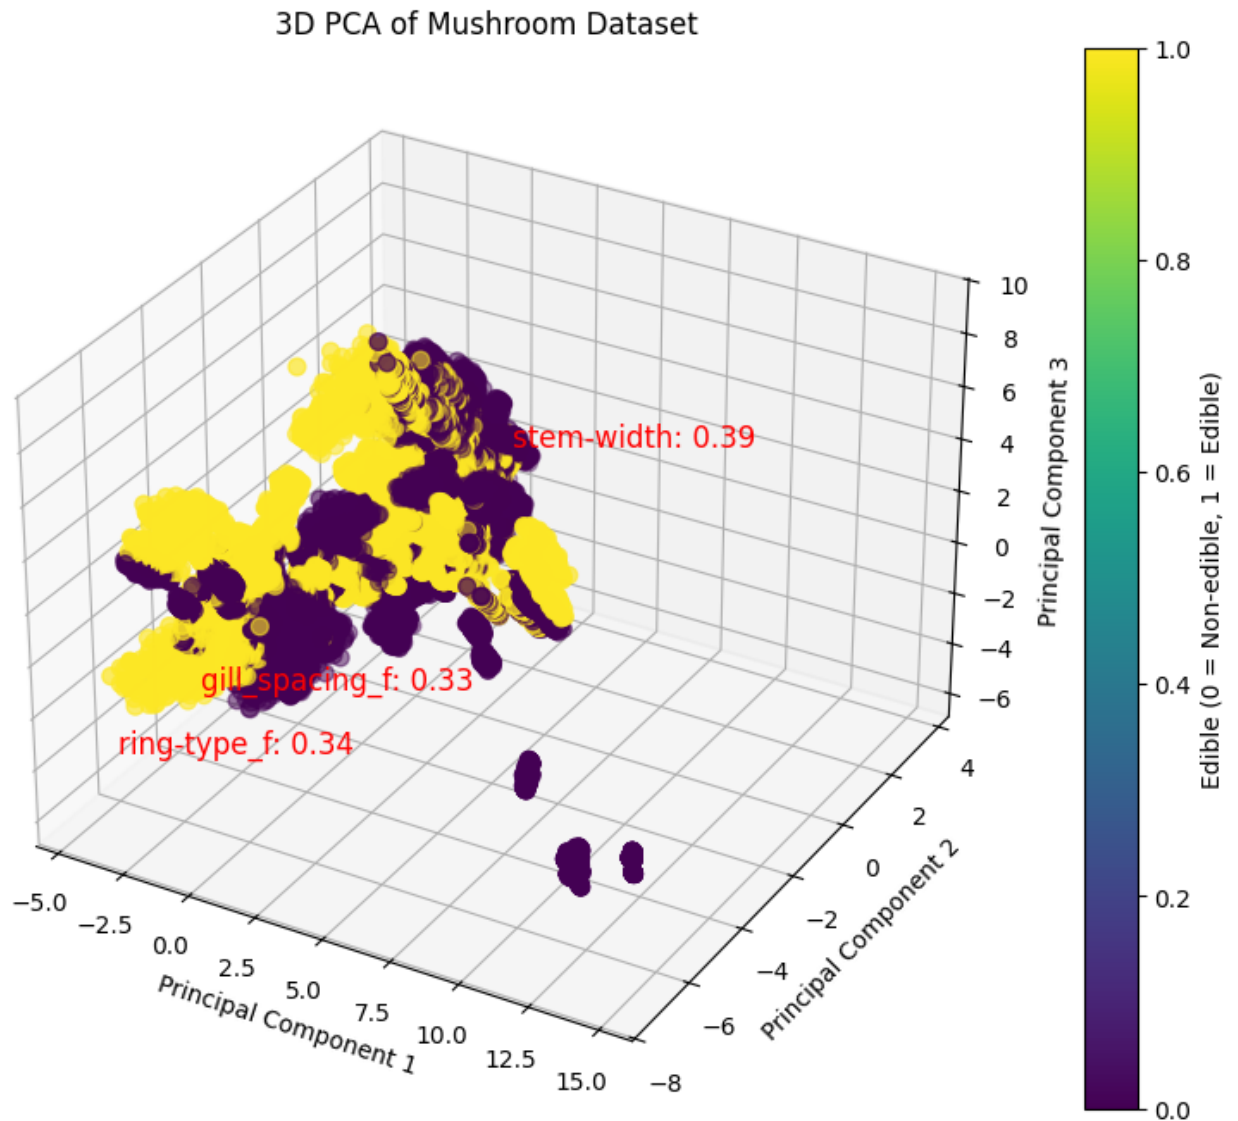
pca_df = pd.DataFrame(X_pca, columns=['PC1', 'PC2', 'PC3'])
pca_df['edible'] = y
components_df = pd.DataFrame(pca.components_, columns=X.columns, index=['PC1', 'PC2', 'PC3'])

# Plot the 3D PCA plot
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
ax.set_xlabel('Principal Component 1')
ax.set_ylabel('Principal Component 2')
ax.set_zlabel('Principal Component 3')
ax.set_title('3D PCA of Mushroom Dataset')

# Color points based on the 'edible' target (0 for non-edible, 1 for edible)
scatter = ax.scatter(pca_df['PC1'], pca_df['PC2'], pca_df['PC3'], c=pca_df['edible'])

fig.colorbar(scatter, label='Edible (0 = Non-edible, 1 = Edible)')
for i in range(3): # Loop through the 3 components
    ax.text(
        pca_df.iloc[0, i], pca_df.iloc[1, i], pca_df.iloc[2, i],
        f'{components_df.iloc[i].idxmax(): {components_df.iloc[i].max():.2f}}',
        color='red', fontsize=12
    )

plt.show()
```



The 3D PCA shown above shows some grouping that is clear in their inedibility due to stem\_width, we saw this earlier in the feature importance, so this is supporting that the model will be able to classify with accuracy. While there are still groupings that can lead to easy classification by the 3 components, they are not as clear and separated.

### Backpropogated Perceptron Neural Netowrk

For this model, we will be using a 3 layer backpropogated perceptron neural network. The activation function will be sigmoid to as ReLu could drop smaller features from contributing to the classification.

```
In [11]: # seperate features and target
X = df1.drop('edible', axis=1)
Y = df1['edible']
# spit data into training and testing sets 80:20
```

```

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, ran

#reshape to vertical vector
Y_train = Y_train.values.reshape(-1, 1)
Y_test = Y_test.values.reshape(-1, 1)

# print splits
print("Training feature set shape:", X_train.shape)
print("Testing feature set shape:", X_test.shape)
print("Training labels shape:", Y_train.shape)
print("Testing labels shape:", Y_test.shape)

```

Training feature set shape: (48738, 116)

Testing feature set shape: (12185, 116)

Training labels shape: (48738, 1)

Testing labels shape: (12185, 1)

With the data now split, we need to build and train the model.

```

In [18]: # Arrays to store the loss values
train_losses = []
test_losses = []

# layer perceptron neuron counts
input_neurons = X.shape[1]
hidden_neurons_1 = 21 # first hidden layer
hidden_neurons_2 = 7 # second hidden layer
output_neurons = 1

# initialize random weights and biases
np.random.seed(42)
W1 = np.random.randn(input_neurons, hidden_neurons_1) * np.sqrt(2 / input_ne
b1 = np.zeros((1, hidden_neurons_1))
W2 = np.random.randn(hidden_neurons_1, hidden_neurons_2) * np.sqrt(2 / hidde
b2 = np.zeros((1, hidden_neurons_2))
W3 = np.random.randn(hidden_neurons_2, output_neurons) * np.sqrt(2 / hidden_
b3 = np.zeros((1, output_neurons))

# activation function
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    return x * (1 - x)

# Hyperparameters
learning_rate = 0.05
epochs = 30000

# Training loop

```



```
for epoch in range(epochs):
    # Forward pass
    Z1 = np.dot(X_train, W1) + b1
    A1 = sigmoid(Z1)
    Z2 = np.dot(A1, W2) + b2
    A2 = sigmoid(Z2)
    Z3 = np.dot(A2, W3) + b3
    A3 = sigmoid(Z3)

    # Compute cross-entropy loss
    epsilon = 1e-10 # To avoid log(0)
    loss = -np.mean(Y_train * np.log(A3 + epsilon) + (1 - Y_train) * np.log(1 - A3 + epsilon))
    train_losses.append(loss)

    # Backward pass
    dZ3 = A3 - Y_train
    dW3 = np.dot(A2.T, dZ3) / X_train.shape[0]
    db3 = np.sum(dZ3, axis=0, keepdims=True) / X_train.shape[0]

    dA2 = np.dot(dZ3, W3.T)
    dZ2 = dA2 * sigmoid_derivative(A2)
    dW2 = np.dot(A1.T, dZ2) / X_train.shape[0]
    db2 = np.sum(dZ2, axis=0, keepdims=True) / X_train.shape[0]

    dA1 = np.dot(dZ2, W2.T)
    dZ1 = dA1 * sigmoid_derivative(A1)
    dW1 = np.dot(X_train.T, dZ1) / X_train.shape[0]
    db1 = np.sum(dZ1, axis=0, keepdims=True) / X_train.shape[0]

    # Update weights and biases
    W3 -= learning_rate * dW3
    b3 -= learning_rate * db3
    W2 -= learning_rate * dW2
    b2 -= learning_rate * db2
    W1 -= learning_rate * dW1
    b1 -= learning_rate * db1

    # Forward pass (testing data)
    Z1_test = np.dot(X_test, W1) + b1
    A1_test = sigmoid(Z1_test)
    Z2_test = np.dot(A1_test, W2) + b2
    A2_test = sigmoid(Z2_test)
    Z3_test = np.dot(A2_test, W3) + b3
    A3_test = sigmoid(Z3_test)

    # Compute testing loss (cross-entropy)
    test_loss = -np.mean(Y_test * np.log(A3_test + epsilon) + (1 - Y_test) * np.log(1 - A3_test + epsilon))
    test_losses.append(test_loss)

    # print testing loss
```

```
if epoch % 1000 == 0:  
    print(f"Epoch {epoch}, Loss: {loss:.4f}")
```

```
Epoch 0, Loss: 0.7026  
Epoch 1000, Loss: 0.6829  
Epoch 2000, Loss: 0.6649  
Epoch 3000, Loss: 0.6043  
Epoch 4000, Loss: 0.5227  
Epoch 5000, Loss: 0.4597  
Epoch 6000, Loss: 0.4170  
Epoch 7000, Loss: 0.3822  
Epoch 8000, Loss: 0.3495  
Epoch 9000, Loss: 0.3172  
Epoch 10000, Loss: 0.2845  
Epoch 11000, Loss: 0.2526  
Epoch 12000, Loss: 0.2226  
Epoch 13000, Loss: 0.1938  
Epoch 14000, Loss: 0.1655  
Epoch 15000, Loss: 0.1384  
Epoch 16000, Loss: 0.1139  
Epoch 17000, Loss: 0.0929  
Epoch 18000, Loss: 0.0757  
Epoch 19000, Loss: 0.0621  
Epoch 20000, Loss: 0.0513  
Epoch 21000, Loss: 0.0429  
Epoch 22000, Loss: 0.0363  
Epoch 23000, Loss: 0.0311  
Epoch 24000, Loss: 0.0270  
Epoch 25000, Loss: 0.0236  
Epoch 26000, Loss: 0.0208  
Epoch 27000, Loss: 0.0186  
Epoch 28000, Loss: 0.0167  
Epoch 29000, Loss: 0.0151
```

Next we will test the accuracy of the model using the previously separated testing data.

\*Note: While the hyperparameters were modified in order to try and achieve a better performing model, using LOTS of epochs proved to train the model best and denote the best performing model as possible.

```
In [22]: # Classify accuracy using test data  
Z1_test = np.dot(X_test, W1) + b1  
A1_test = sigmoid(Z1_test)  
Z2_test = np.dot(A1_test, W2) + b2  
A2_test = sigmoid(Z2_test)  
Z3_test = np.dot(A2_test, W3) + b3  
A3_test = sigmoid(Z3_test)  
  
# Convert predictions to binary
```

```
Y_pred = (A3_test > 0.5).astype(int)

# Print accuracy and classification report
accuracy = accuracy_score(Y_test, Y_pred)
print(f"Test Accuracy: {accuracy:.4f}")
print("\nClassification Report:")
print(classification_report(Y_test, Y_pred))
```

Test Accuracy: 0.9990

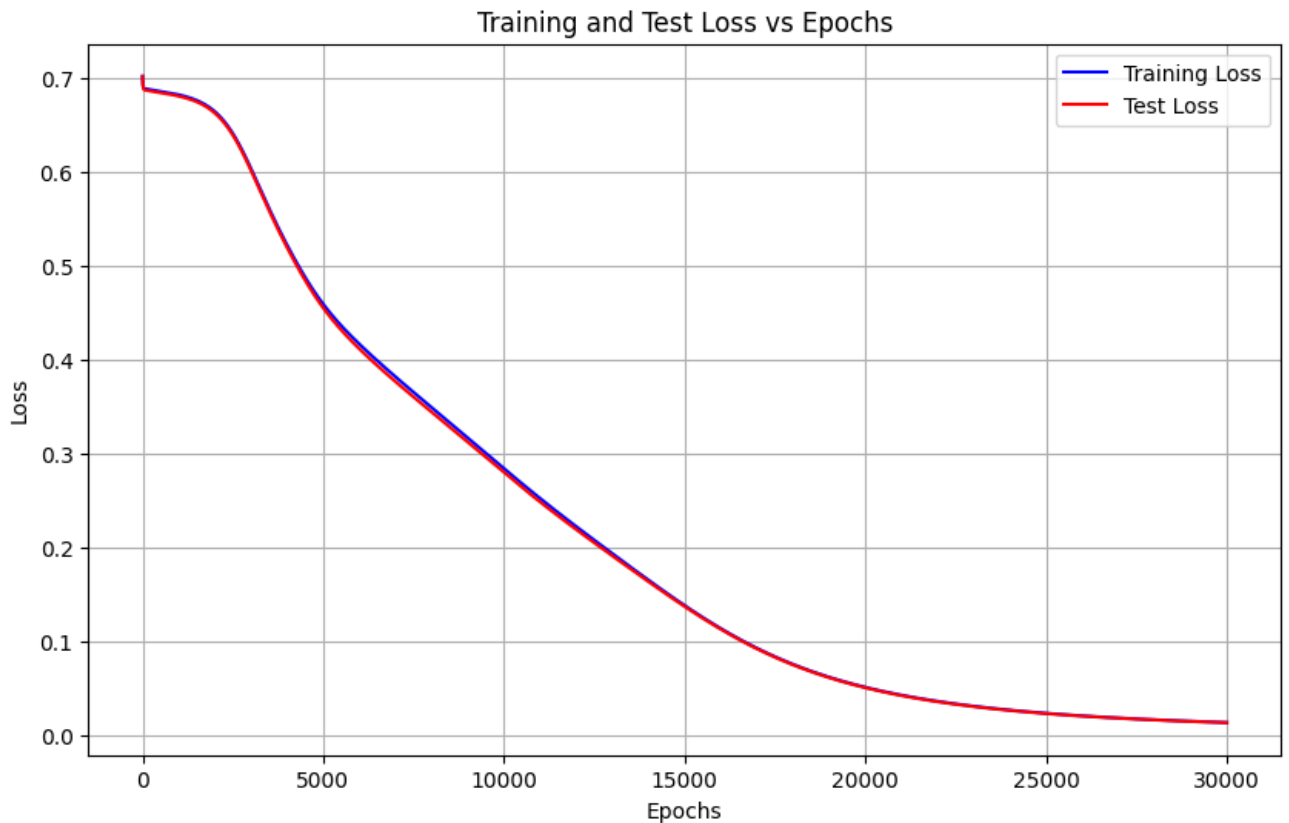
Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6820
1	1.00	1.00	1.00	5365
accuracy			1.00	12185
macro avg	1.00	1.00	1.00	12185
weighted avg	1.00	1.00	1.00	12185

From the accuracy and classification report, we can see that the model is performing very well as classifying if a fungi is edible or posinous.

The precision for both the outcomes are 100%, meaning the model is able to always correctly identify if a mushroom is edible or poisonous based on its physical characteristics.

```
In [23]: # Plotting the losses over epochs to check for overfitting
plt.figure(figsize=(10, 6))
plt.plot(range(epochs), train_losses, label='Training Loss', color='blue')
plt.plot(range(epochs), test_losses, label='Test Loss', color='red')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Test Loss vs Epochs')
plt.legend()
plt.grid(True)
plt.show()
```



As we can see from the graph above, the training loss and test data loss consistently gets smaller. If at some the training loss were to continue to reduce but the test loss were to begin increasing that would denote overfitting. While I was worried that the number of epochs that were used to get such a high classification score would lead to over fitting, it appears to not have and is just a well trained model.

The test below is again to help us determine if the model is overfitting by checking if the test data loss is ever more than the training loss. Since the test loss is never more than the training loss, it shows us that there is no overfitting and the model is able to correctly classify new data.

```
In [24]: # Check if the test loss is ever greater than the training loss
epochs_with_higher_test_loss = [epoch for epoch, (train_loss, test_loss) in

if epochs_with_higher_test_loss:
    print(f"The test loss exceeds the training loss at the following epochs:
else:
    print("The test loss is never greater than the training loss.")
```

The test loss is never greater than the training loss.